

ECS 170: Learning to Play Pong

Ethan He

1 Problem Representation

1. Since we are using a neural network to replace the Q-table, it is easier to use image as state input. As it did in the `QLearner` class:

```
1 self.features = nn.Sequential(  
2     nn.Conv2d(self.input_shape[0], 32, kernel_size=8, stride=4),  
3     nn.ReLU(),  
4     nn.Conv2d(32, 64, kernel_size=4, stride=2),  
5     nn.ReLU(),  
6     nn.Conv2d(64, 64, kernel_size=3, stride=1),  
7     nn.ReLU()  
8 )  
9  
10 self.fc = nn.Sequential(  
11     nn.Linear(self.feature_size(), 512),  
12     nn.ReLU(),  
13     nn.Linear(512, self.num_actions)  
14 )
```

2. The purpose of the neural network in Q-Learning is to replace the lookup table. In such way, we can train a network for each action, where we use state as input to the network and get \hat{Q} as output.
3. ϵ is the identifier for choosing an action. We generate a random number, if it is greater than ϵ , we do exploitation, so choose the best known action that the Q learner tell us; otherwise, we do exploration, then do random action.
4. See function `act` of `dpn.py`.

2 Making a Q-Learner Learn

The loss function is the square error formula. In the function, we get the current environment, where `batch_size` controls the sample size we fetch from `replay_buffer` (tensor). And from the random samples, we find the source to calculate the square error. The parameter `gamma` is to control the Q-learner's behavior: as γ get closer to 1, future rewards are given greater emphasis relative to the immediate rewards.

```
1 q_val = torch.gather(model(state.squeeze(1)), 1, action.unsqueeze(1)).  
    squeeze(1) # y_i  
2 q_val_target = target_model(next_state.squeeze(1)).max(1)[0]  
3 expectations = reward + gamma * q_val_target * (1 - done)  
4 loss = (q_val - Variable(expectations.data)).pow(2).mean()
```

3 Extend the Deep Q-Learner

See program dqn.py.

```
1 def sample(self, batch_size):
2     # Randomly sampling data with specific batch size from the buffer
3     state, action, reward, next_state, done = zip(*random.sample(self.
4     buffer, batch_size))
5
6     state = Variable(torch.FloatTensor(np.float32(state)))
7     next_state = Variable(torch.FloatTensor(np.float32(next_state)))
8
9     return state, action, reward, next_state, done
```

4 Learning to Play Pong

1. See program run_dqn_pong.py around line 90.


```
1 torch.save(model.state_dict(), model_file)
```

2. Use the np.savetxt() function to save losses and all_rewards in .csv files for use in question 4.


```
1 np.savetxt("rewards.csv", all_rewards, delimiter=",")
2 np.savetxt("losses.csv", losses, delimiter=",")
```

3. Use command `python3 run_dqn_pong.py` to train the model.

4. Plots.



(a) Rewards



(b) Losses

5 Bonus